

Feasibility of Homomorphic Inference for a Genomic Foundation Model: A Client-Assisted CKKS Evaluation of DNAGPT

Christos Galanopoulos^{1,*}, Kimon Antonios Provatas¹, Ilias Georgakopoulos-Soares^{1,*}

¹The University of Texas at Austin, Austin, TX, USA

*Correspondence: Christos Galanopoulos and Ilias Georgakopoulos-Soares

Homomorphic encryption · Genomic privacy · Transformer inference · GPU acceleration

Abstract

Motivation. Genomic sequence is identifying and cannot be replaced after disclosure. We ask whether a provider can execute a released genomic transformer without receiving the genomic input in plaintext, and where numerical agreement, memory, or cost limits that deployment.

Results. We validate DNAGPT on three genomic task families and construct an independently checked numerical reference. In the client-assisted design, the provider evaluates every linear operation on ciphertexts under one key lineage; the key holder evaluates LayerNorm inverse square roots, causal softmax, and GELU exactly, at boundaries fixed before execution. All 12 released blocks and the task head execute from encrypted embedded vectors with declared client refreshes and no homomorphic bootstrap, leaving the released weights and nonlinear formulas unmodified. For the one prompt evaluated, the decrypted label matches the reference and its margin has 8.56×10^{-9} relative error, 6.67 orders of magnitude below the 4×10^{-2} tolerance fixed in advance; no polynomial approximation of a nonlinearity contributes to that residual. Client refresh holds multiplicative depth at 13, and device-wide GPU memory used reaches 9839 MiB without growing by block index. One in-process execution takes 1.86 hours on an A100 GPU node. A pure non-interactive configuration completes one released-weight block, then exceeds the tested 80 GB accelerator-memory envelope on its composition context and evaluation keys, placing that barrier in the nonlinear schedule rather than in the model’s dense maps.

Conclusions. The executed case establishes complete-model graph feasibility and numerical agreement for one input, and locates the limiting barrier in throughput rather than in correctness or accelerator memory. Client refresh bounds multiplicative depth by the longest encrypted segment rather than by network depth, which is what makes the twelve-block lineage expressible without bootstrapping; stage-bounded plaintext caching separately prevents host-side accumulation. The current single-process implementation is not practical for interactive use and has not been evaluated as a networked service, so feasibility and practicality are reported as separate verdicts.

Availability. Code, evaluation harnesses, and the evidence record are maintained at https://github.com/Georgakopoulos-Soares-lab/private_genomic_dnagpt.

Key Points

- All twelve released DNAGPT blocks and the task head execute from encrypted embedded vectors at the 103-token task prompt, under one key lineage and without any homomorphic bootstrap. The released weights and nonlinear formulas are preserved unmodified, and the one evaluated prompt reproduces the reference label with 8.56×10^{-9} relative margin error.
- The validation chain separately records released float32 behavior, checks an independent NumPy float64 reference against a float64-cast PyTorch shadow, and compares encrypted execution with that reference.
- Because client refresh bounds depth by the longest encrypted segment rather than by network depth, multiplicative depth stays at 13 and device-wide GPU memory used stays at 9839 MiB across the complete model, with no growth by block index. Stage-bounded caching separately limits host memory while preserving the asserted operation schedule.
- One in-process execution takes 1.86 hours and requires an online key holder, so the evaluated form is not interactive. Feasibility and practicality are therefore reported as separate verdicts: complete-model feasibility is established for the evaluated prompt, while encrypted task accuracy, service rate, and networked deployment remain open.

A genomic fragment stays off the compute provider




Figure 1: The data owner tokenizes and embeds a genomic fragment locally, encrypts the embedded vectors, and retains the secret key. The provider evaluates model linear algebra on ciphertexts; the data owner evaluates declared nonlinear functions on CPU. The measured implementation emulates both roles in one process. The result panel reports one measured execution at the 103-token genomic-signal prompt.

1 Introduction

Human genomic sequence carries an unusually persistent confidentiality risk. It can identify its contributor and, through familial relationships, reveal information about people who never supplied a sample; re-identification has been demonstrated even when conventional identifiers were removed [Gymrek et al. \[2013\]](#), [Erlich et al. \[2018\]](#). A disclosed genome also cannot be replaced in the manner of a password or access token. European data-protection law accordingly treats genetic data as a special category [European Parliament and Council of the European Union \[2016\]](#). The same sequence is the input on which genomic foundation models derive their value: DNAGPT serves classification, regression, and generation tasks through a common transformer backbone [Zhang et al. \[2023\]](#). When that model is served rather than distributed, a data owner without accelerator hardware must otherwise disclose the sequence to the provider or forgo the inference. Nor is handing over an embedding a way out: per-token embeddings from DNA foundation models have been shown to permit near-perfect reconstruction of the underlying sequence [Ouaari et al. \[2026\]](#). Homomorphic encryption offers a route for the provider to compute without receiving the genomic values in plaintext.

DNAGPT’s public release makes the experiment auditable and also makes this particular artifact available for local execution. We use it as a reproducible surrogate for a genomic model offered only as a service by policy or contract. Under that premise, client-assisted evaluation gives a CPU-equipped data owner access to the served model without disclosing the genome.

Against this deployment premise, we ask: can inference with a released genomic foundation model be executed under homomorphic encryption, and where would correctness, performance, or memory prevent a complete encrypted deployment? The question contains two verdicts that should not be collapsed. Feasibility asks whether the released computation can be executed faithfully within the available cryptographic and memory budget. Practicality asks whether its resource cost supports a target use. A correct but slow encrypted path is therefore evidence: it resolves the first question while locating the work needed to change the second.

The technical difficulty is concentrated where a transformer departs from affine arithmetic. CKKS supports approx-

imate packed arithmetic and maps naturally to dense projections, rotations, and accumulations [Cheon et al. \[2017\]](#); LayerNorm, softmax, and GELU are not native CKKS operations. A non-interactive evaluator must approximate them with polynomial circuits, lengthening the multiplicative path and increasing the context, bootstrapping, and evaluation-key material that must be resident. Moving the exact functions to the key holder shortens those paths and resets the depth budget, but introduces interaction and a different threat model. Earlier systems establish both directions: THOR demonstrates non-interactive homomorphic transformer inference; THE-X and Saffire use interactive workflows in which the client evaluates nonlinear functions [Moon et al. \[2025\]](#), [Chen et al. \[2022\]](#), [Biswas et al. \[2025\]](#). The division of labor is therefore established prior art. What has not been reported is an auditable account of the numerical behavior, memory, interaction, and executed work that results when it is applied end to end to a released genomic transformer, at that model’s own task prompt, with its weights and nonlinear formulas unmodified.

That choice buys exactness and depth. Evaluating LayerNorm, softmax, and GELU at the key holder keeps the released formulas intact, so no polynomial approximation of a nonlinearity enters the model, and it bounds multiplicative depth by the longest encrypted segment rather than by network depth. The price is an online key holder and a protocol that is not the fastest available: non-interactive systems report lower cost for transformer inference, and Section 8 treats those figures as cross-paper context rather than as a controlled benchmark.

We begin by establishing that the target computation is meaningful and independently reproducible. The released DNAGPT paths are evaluated across genomic-signal recognition, promoter and splice-site classification, and mRNA abundance regression, with task outputs frozen before encrypted work. The validation records the released float32 path, checks an independent NumPy float64 lineage against a float64-cast PyTorch shadow at 2×10^{-5} for every released block and the task head, and then compares encrypted execution with the NumPy reference. We build two CKKS designs: a non-interactive circuit with polynomial nonlinearities, and a client-assisted circuit whose fixed, non-adaptive boundaries evaluate the released nonlinearities exactly. Instrumented evaluations record where each design stops in correctness and memory, while runtime assertions preserve the complete cryptographic operation schedule throughout the systems campaign.

The results separate three candidate barriers and resolve two of them. Correctness holds: the client-assisted design evaluates all twelve released blocks and the task head at the 103-token prompt, and for the one input evaluated it reproduces the frozen plaintext label with 8.56×10^{-9} relative margin error against a 4×10^{-2} tolerance fixed before evaluation. Accelerator memory holds as well, and for a structural reason: because multiplicative depth is bounded by the longest segment between client boundaries instead of by network depth, it stays at 13 for the complete model, device-wide GPU memory used reaches 9839 MiB without growth across the twelve blocks, and no homomorphic bootstrap is required at any point. Bounded caching separately prevents encoded plaintexts from accumulating across host-side stages. The evaluated non-interactive design, by contrast, completes one released-weight block but exceeds the tested 80 GB envelope while loading its composition context and evaluation keys, placing that barrier in the nonlinear schedule rather than in the model’s dense maps. Cost is the barrier that remains: one encrypted classification takes 1.86 hours on an A100 GPU node, which forecloses interactive use in the evaluated form, and sampled GPU utilization peaks at 51%, locating headroom without attributing it. These memory and timing figures are specific to the evaluated backend, packing, and parameters.

The paper makes the following contributions:

- **Complete-model feasibility.** All twelve released DNAGPT blocks and the task head execute at the 103-token task prompt from encrypted embedded vectors, with unmodified released weights and nonlinear formulas and without homomorphic bootstrap; the one evaluated prompt reproduces the reference label. To our knowledge this is the first complete forward pass of a released, pretrained genomic foundation model evaluated under CKKS; prior homomorphic transformer evaluations over biological sequence data have targeted small, purpose-trained classifiers [Harmanci et al. \[2024\]](#), [Yang et al. \[2024\]](#).
- **Model-faithful, reproducible evaluation.** The protocol, packing, checked operation schedule, predeclared tolerance, and three-link float32–float64–encrypted validation chain make the complete execution independently inspectable.
- **Measured systems anatomy and scoped negative results.** We report complete-model cost, interaction, depth, and memory; the stage-bounded cache that makes host-side reuse executable; and the memory boundary reached by the evaluated non-interactive configuration.

2 Background

2.1 DNAGPT and the genomic task families

DNAGPT is an autoregressive transformer adapted to sequence and numerical tasks over DNA [Zhang et al. \[2023\]](#). We evaluate its released 0.1-billion-parameter configuration. The backbone has 12 transformer blocks, hidden width

768, and 12 causal-attention heads of width 64. Its feed-forward sublayer expands to width 3,072. DNA is represented through 6-mer tokenization and combined with task-format tokens before entering the transformer.

Each block follows the structure that determines the encrypted circuit: LayerNorm; query, key, and value projections; causal self-attention; an output projection and residual; a second LayerNorm; then an expanded MLP with GELU and another residual. The dense maps use fixed model weights, whereas LayerNorm, attention, and GELU introduce data-dependent products or nonlinear functions.

The release includes fine-tuned heads for genomic-signal recognition and mRNA abundance regression. It does not include a head for the genome-understanding benchmark, so Section 4 evaluates that backbone with a locally fine-tuned linear head. Each of these families predates the model: signal recognition by task-specific convolutional architectures Kalkatawi et al. [2019], abundance regression likewise Agarwal and Shendure [2020], and the understanding benchmark by the encoder-based model that introduced it Zhou et al. [2024]. They first establish model fidelity; the encrypted study then uses the signal-recognition path as its task-length anchor.

2.2 CKKS

CKKS supports approximate arithmetic on packed vectors of real or complex values Cheon et al. [2017]. Encoding places many values into the slots of one plaintext, which is then encrypted under the data owner’s key. Addition and multiplication act slotwise across the packed vector, while rotations move values between slots. This single-instruction, multiple-data model is a natural match for transformer activations: a rotation-and-diagonal schedule can express a public matrix transform as ciphertext–plaintext multiplication and accumulation.

Not every transformer product has a public operand. Query–key scoring multiplies two encrypted projections of the input, and attention context multiplies encrypted attention weights by encrypted values. Those operations require ciphertext–ciphertext multiplication and relinearization. The distinction is about data ownership rather than implementation convenience: public weights and masks remain plaintext, while every query-derived operand remains encrypted at the server.

CKKS controls approximate fixed-point scale through rescaling. A multiplication consumes a level from a finite modulus chain, and the computation fails once that multiplicative budget is exhausted. Bootstrapping can refresh a ciphertext without decryption, while fresh client encryption can reset the budget in an interactive protocol. Both choices carry system costs, so multiplicative depth is an architectural resource rather than only a cryptographic parameter.

2.3 Why the nonlinearities are the architecture

CKKS directly supplies polynomial arithmetic, but a transformer requires inverse square root in LayerNorm, exponentiation and normalization in softmax, and GELU. A non-interactive circuit must replace these functions with polynomial approximations over declared input ranges. Higher-degree approximations consume additional multiplicative depth; across repeated blocks, the evaluator must provide a deeper context, a bootstrapping schedule, or both. Larger modulus chains and refresh operations also increase the context and evaluation-key material that must remain resident.

The alternative is to place exact functions at declared client boundaries. The key holder decrypts only values derived from its query, evaluates the nonlinearity, and returns a fresh ciphertext. This removes approximation-range calibration and resets the encrypted depth budget, but it requires an online client and exposes intermediate activations to that client. The choice between polynomial approximation and client evaluation therefore determines memory, interaction, depth, and the threat model. Section 5 evaluates both designs and reports where each stops.

2.4 Evaluated software backends

We separate correctness semantics from the performance path. OpenFHE provides CKKS context and security-parameter construction, encoding, rotations, polynomial evaluation, and approximate bootstrapping Al Badawi et al. [2022]. It supplies the reference behavior used to validate local encrypted operators and the non-interactive design. The evaluated GPU backend interoperates with those representations and provides CUDA implementations of the server-side CKKS operations; all reported accelerator measurements use that path. Python, NumPy, and PyTorch construct and validate the frozen plaintext reference, but they do not participate in encrypted performance measurements.

We also evaluated the Boolean/integer-oriented TFHE-rs path and Concrete ML’s client/server workflow Zama [2022a,b]. In this project’s backend comparison, their available GPU acceleration and ciphertext expansion were less suitable for this matrix-heavy graph than the already validated CKKS path. A separate CKKS wrapper was not selected because its public interface lacked the refresh and general nonlinear-function support required by the

Table 1: Resource and trust allocation under local and served-model deployment. The plaintext service is the ordinary baseline that the encrypted protocol is intended to change.

	Client-local	Plaintext service	Client-assisted CKKS
Client must hold the full model	yes	no	no
Server receives genomic input	no	in plaintext	encrypted vectors
Client requires a GPU	deployment-dependent	no	no
Server evaluates model linear algebra	no	in plaintext	on ciphertexts
Client evaluates model nonlinearities	locally	no	at boundaries

non-interactive baseline. These decisions select among the software paths tested in this project. They do not establish that CKKS is the only scheme able to express encrypted transformer inference.

3 Deployment scenario and threat model

3.1 Deployment roles and privacy objective

The deployment has two parties but one cryptographic confidentiality objective. The data owner holds a human genomic sequence, ordinary CPU resources, and no GPU. It will not disclose the sequence to the compute provider. The model owner operates the accelerator infrastructure and serves the DNAGPT model rather than distributing it to clients. Keeping the weights server-side is an operational preference of this deployment, not a model-confidentiality guarantee supplied by the protocol.

The released DNAGPT artifact makes the experiment reproducible; it does not itself require this service arrangement. A client may run the public release locally. We use it as an auditable surrogate for a genomic model whose owner limits access to a service by policy or contract, and the deployment conclusions apply under that premise.

The protected asset is the genomic input. The data owner encrypts its embedded numeric vectors and retains the secret key; the compute provider evaluates the served model on those ciphertexts. The protocol is relevant only where these roles are fixed. If a client is allowed to receive and run the model, local plaintext inference is the simpler architecture.

The evaluated input path assumes that the provider distributes the tokenizer and embedding table so that the data owner can construct embedded vectors locally. This is compatible with the public artifact used in the experiment, but it is an explicit service assumption. A provider that will not distribute those components would instead need a private token-index lookup, which is outside the evaluated protocol.

3.2 Local inference as a comparison

Local plaintext inference is preferable whenever the client may hold and run the model. It avoids all cryptographic work, and client assistance offers no computational advantage over it. In the complete encrypted inference reported in Section 7, the data owner spends 956.187 s evaluating LayerNorm inverse square roots, causal softmax, and GELU at the boundaries. This study records no comparable local CPU forward-pass time, so it does not quantify the overhead. The evaluated weights are also public, and no model-confidentiality argument requires the service arrangement.

The protocol instead addresses a served-model deployment in which the client may not hold the weights. Under that premise, the choices are to disclose the genome, decline the inference, or use a privacy-preserving protocol. Client-assisted CKKS provides access to the served model without genomic disclosure; it is not intended to outperform plaintext inference.

The protocol assigns encrypted dense linear algebra to the party with accelerator hardware while the key holder evaluates exact nonlinearities on CPU. This allocation protects the genomic input without requiring the data owner to operate a GPU.

This allocation becomes more server-heavy as model width grows. At fixed sequence structure, the client’s boundary work scales with the number of activation values, whereas the server’s dense transforms scale quadratically with hidden width. The client fraction therefore decreases as $O(1/D)$ with hidden width D under this fixed sequence structure. This is an asymptotic property of the work allocation, not a measurement of a wider DNAGPT model.

3.3 Adversary model

We consider a semi-honest compute provider under 128-bit classical CKKS parameters. It follows the prescribed circuit but may inspect everything it legitimately receives: public and evaluation keys, model parameters, ciphertexts,

sequence length, message sizes, operation order, boundary schedule, and timing. It does not alter ciphertexts to probe the client or deviate from the declared boundary sequence.

Under this adversary, the compute provider receives no plaintext genomic vector, plaintext activation, partial decryption, or secret-key material. At a boundary it supplies a ciphertext and receives a fresh ciphertext under the data owner’s public key. The guarantee concerns the provider’s view of query-derived values; public model structure and plaintext weights remain available to the provider as inputs to ciphertext–plaintext operations.

Invariant 1 (Key custody). *The data owner generates and retains the secret key. No protocol message contains the key or a partial decryption, and the compute provider has no decryption operation.*

Invariant 2 (No plaintext at the server). *Every query-derived value held by the compute provider is a ciphertext under the data owner’s key. Client-evaluated nonlinearities return through fresh encryption rather than as decoded values.*

Invariant 3 (Non-adaptive boundaries). *The circuit and sequence length determine every boundary before execution; decrypted content does not select the next operation. The implementation asserts the complete schedule and boundary counts for each accepted run and aborts on deviation. Transcript equality across multiple private inputs of equal length has not been measured as a separate leakage experiment.*

3.4 What is not claimed

Model confidentiality is not claimed. The data owner observes full intermediate activations at every nonlinearity. This transcript creates a chosen-query extraction surface, and sufficiently diverse input–output pairs may make some affine segments identifiable. No extraction experiment or model-confidentiality proof is provided. Keeping weights on the server may deter casual copying, but it is not a cryptographic confidentiality guarantee. This limitation does not change the genomic-data objective, which protects the data owner from the compute provider.

Malicious servers, authenticated transport, production key custody, traffic analysis, timing and physical side channels, and compromised clients are outside the evaluated boundary. Encrypted scope begins after token-index lookup, so private lookup is also unresolved. Per-token embeddings remain sensitive: high-fidelity sequence reconstruction has been demonstrated for embeddings from other DNA foundation models [Ouaari et al. \[2026\]](#). The embedded vectors are therefore encrypted here rather than sent in clear. The data owner sees its own intermediate activations and final output; privacy from the key holder is not an objective of this protocol.

4 Plaintext reference and model fidelity

An encrypted execution is useful only if the computation it reproduces is itself meaningful and independently checked. Task metrics answer the first question: they establish that the released model and locally trained heads retain the genomic capabilities for which the model is being evaluated. Per-example predictions answer the second. They are recorded before any encrypted work and become the frozen plaintext reference against which encrypted outputs are accepted.

The plaintext harness imports the upstream model implementation and preserves its classification and regression forward paths. It adds batching, metric computation, and result recording without changing model arithmetic. The promoter, splice-site, and abundance tasks use their documented evaluation splits. Genomic-signal recognition is evaluated on the full balanced source corpus, including examples used to train the released head, so its aggregate metric is release and pipeline fidelity evidence rather than held-out performance. Each run records the per-example outputs needed downstream.

The reference is not trusted merely because it is written in NumPy. The numerical chain has three explicit links. First, the released model’s default float32 path is recorded as a diagnostic lineage. Second, a deep-copied float64 PyTorch shadow is compared with an independent NumPy float64 re-derivation for every one of the 12 blocks and the classifier head; generation aborts unless relative and absolute differences satisfy 2×10^{-5} . Their largest per-block absolute difference is 5.22×10^{-6} . Third, encrypted execution is compared with the accepted NumPy float64 reference. The released float32 path is not subject to the float64-shadow tolerance: its largest recorded absolute block drift from the NumPy lineage is 7.997×10^{-4} . For the selected full prompt, the released float32, manual float64, and encrypted lineages nevertheless all produce label N; the encrypted margin differs from the released float32 margin by 1.61×10^{-6} in relative terms.

Table 2 shows that the target carries useful signal across classification and regression tasks. On the full polyadenylation-signal source corpus of 22,604 examples, the released classifier reaches 0.9124 accuracy and 0.916 F1, against the 0.9151 accuracy [Zhang et al. \[2023\]](#) report for this model class on a held-out quarter of that set. The locally fine-tuned understanding-benchmark heads reach MCC 0.680, 0.897, and 0.831 on core-promoter, promoter, and splice-site

Table 2: Released DNAGPT reproduced locally on three genomic task families, against published references. Signal recognition and abundance regression use the released fine-tuned heads; the understanding-benchmark rows use a locally fine-tuned head, as no head was released.

Task	Metric	This work	Reference	Source
Polyadenylation-signal recognition [‡]	accuracy	0.9124	0.9151	Zhang et al. [2023]
	F1	0.916	—	
Core promoter detection	MCC	0.680	0.69	Zhou et al. [2024]
300 bp promoter detection	MCC	0.897	0.87	Zhou et al. [2024]
Splice-site classification	MCC	0.831	0.85	Zhou et al. [2024]
mRNA abundance regression ^{†§}	r^2	0.562	0.62	Zhang et al. [2023]
	Pearson r	0.753	—	

[‡]The reference is DNAGPT’s own reported accuracy for this model class, measured on a held-out quarter of the set; this work evaluates all 22,604 examples. The comparison is same-model but not same-split.

[†]Model-fidelity evidence only. This task is outside the encrypted scope; see Section 9.

[§]The released regression head evaluated here is fine-tuned from the human-pretrained 0.1B variant, while the reference is reported for the mammal-pretrained variant, so the backbones differ.

detection, against the 0.69, 0.87, and 0.85 that Zhou et al. [2024] report for DNABERT-2 itself on the same splits; other models in that benchmark score higher on two of the three. The released abundance-regression head reaches $r^2 = 0.562$ and Pearson $r = 0.753$, against the 0.62 Zhang et al. [2023] report over the Agarwal and Shendure [2020] regression baseline. These results establish model fidelity; encrypted evaluation is still required to establish cryptographic feasibility. The genomic-signal aggregate is not presented as an independent generalization result because the full source corpus includes head-training examples.

Two dataset qualifications belong with these baseline results. DNAGPT does not release a head for the understanding benchmark, so we trained a linear head and evaluated the three human promoter and splice-site datasets selected for this study. For abundance regression, the historical source was recovered and its split procedure reproduced, but the resulting test set may not be gene-for-gene identical to the historical split. The abundance result is therefore model-fidelity evidence and is not used as an encrypted target.

4.1 Why 103 tokens

The encrypted anchor length follows from the genomic-signal task rather than from a convenient cryptographic setting. Its 600-base-pair window becomes 100 sequence tokens under the model’s 6-mer tokenization; three task-format special tokens bring the complete prompt to 103 tokens. The evaluated prompt is therefore the task input consumed by the released classifier, rather than a shortened surrogate chosen to fit the encrypted implementation.

The same arithmetic determines the other classification prompts. Core-promoter detection maps a 70-base-pair window to 12 sequence tokens plus one special token, or 13 total. The 300-base-pair promoter task uses 50 sequence tokens plus one special, or 51; the 400-base-pair splice task uses 67 sequence tokens plus one special, or 68. The 103-token genomic-signal prompt is the longest of these encrypted-scope tasks and consequently provides the anchor for the packing and sequence-length analysis in Sections 5 and 7.

5 Client-assisted CKKS protocol

5.1 Nonlinearity boundaries

CKKS carries the affine maps, reductions, residual additions, and attention products without exposing an activation to the compute provider. The difficulty lies in three nonlinearity classes: the inverse square root in LayerNorm, causal softmax, and GELU. We place a declared client boundary at each class. For LayerNorm, the server computes the encrypted mean and variance and sends only the variance term needed for the inverse square root. For attention, it sends encrypted causal score tiles and receives encrypted softmax weights. The GELU boundary operates on the packed MLP activations.

At every boundary, the data owner decrypts values derived from its own query, evaluates the declared function exactly in double precision, and encrypts the result afresh. The data owner already holds the secret key and is the only party that decrypts. The compute provider receives neither a plaintext activation nor the secret key. Between these declared calls, the ciphertext lineage remains on the server and every model operation is evaluated homomorphically.

Table 3: Notation. All values are those of the evaluated configuration.

Symbol	Value	Meaning
D	768	hidden width of the released model
H	12	attention heads, of $D/H = 64$ dimensions each
D_{mlp}	3072	MLP expansion width, $4D$
T	103	tokenized prompt length for the genomic-signal task
B	8	tokens packed per ciphertext (token lanes)
G	13	ciphertext groups, $\lceil T/B \rceil$
W	1024	feature width padded to a power of two
C	4	activation copies per ciphertext, D_{mlp}/D
N	65536	polynomial ring degree
$N/2$	32768	usable plaintext slots per ciphertext, $C \cdot W \cdot B$
L	13	multiplicative depth
λ	128	classical security level, in bits

5.2 Non-interactive design

The non-interactive design keeps a single ciphertext lineage and replaces every nonlinearity with a polynomial approximation. It completed one released-weight transformer block, establishing that the block arithmetic can be expressed without intermediate decryption. The stopping point appeared when the same design was configured for multi-block composition: context and evaluation-key loading exceeded the tested 80 GB GPU memory envelope before useful composition arithmetic began.

The mechanism matters more than the failure location. Repeated polynomial approximation lengthens the multiplicative path, which requires a deeper modulus budget and a bootstrapping schedule. The associated context and evaluation-key material must be resident before evaluation, so composition failed as a memory problem rather than as a numerical-correctness problem. This finding is specific to the evaluated library, parameters, packing, and hardware; it does not establish a general lower bound for non-interactive CKKS. Published non-interactive systems evaluate encoders of comparable depth and width on a single accelerator at these parameters, as Section 8 records, so this subsection reports where our non-interactive configuration stopped and not where the approach stops. It does identify the measured pressure that the client-assisted design removes: exact client evaluation shortens each encrypted segment and replaces encrypted bootstrapping with fresh encryption at already declared boundaries.

5.3 Protocol algorithms


Fresh encryption in Algorithm 2 returns the activation to level 0, so every client boundary fully resets the multiplicative budget. The configured depth of 13 therefore bounds the longest encrypted run between consecutive boundaries. It does not accumulate across the 12 transformer blocks.

Data owner

holds the genomic fragment and secret key · no GPU

Compute provider

holds the model · encrypted activations only · one A100



Only ciphertexts cross the in-process role boundary; fresh encryption resets level to 0.

Figure 2: Evaluation order for one transformer block and party allocation. The provider computes both LayerNorm statistics under encryption; the data owner evaluates the inverse square root, causal softmax, and GELU at declared boundaries. A validation readout decrypts the output, while composition decrypts and freshly re-encrypts the hidden state before the next block.

Algorithm 1 Evaluating one transformer block. $G = 13$ token groups; $\odot$ is homomorphic multiplication; every operation outside a CLIENT call runs on ciphertexts.

Require: encrypted activations $c_0, \dots, c_{G-1}$; public weights W

Stage 1: normalization and projection

```

1: for  $g \leftarrow 0$  to  $G - 1$  do
2:    $n \leftarrow \text{LAYERNORM}(c_g, \gamma_1)$ 
3:    $\beta \leftarrow \text{BABYROTATIONS}(n)$ 
4:    $Q_g, K_g, V_g \leftarrow \text{MATMUL}(\beta, W_q), \text{MATMUL}(\beta, W_k), \text{MATMUL}(\beta, W_v)$ 
5: end for
6: FLUSHENCODEDWEIGHTS

```

▷ Algorithm 3
▷ $N_1 - 1$ rotations, shared
▷ bounds host memory to one stage

Stage 2: causal attention scores

```

7: for  $k \leftarrow 0$  to  $G - 1$  do
8:    $\tilde{K} \leftarrow \{\text{LANESHIFT}(K_k, \delta)\}$  for each  $\delta$  required by a later query group
9:   for  $q \leftarrow k$  to  $G - 1$  do
10:     $S \leftarrow 0$ 
11:    for  $\delta \in \text{ACTIVESHIFTS}(q, k)$  do
12:       $P \leftarrow Q_q \odot \tilde{K}_\delta$ 
13:      for  $h \leftarrow 0$  to  $H - 1$  do
14:         $S \leftarrow S + \text{REDUCESUM}(P \odot \text{mask}_h) \odot \text{isolate}(q, k, \delta, h) / \sqrt{d_h}$ 
15:      end for
16:    end for
17:    CLIENT.TAKESCORETILE( $S, q, k$ )
18:  end for
19: end for
20: CLIENT.SOFTMAX

```

▷ decrypt only
▷ stable softmax over complete causal rows

Stage 3: attention context

```

21: for  $k \leftarrow 0$  to  $G - 1$  do
22:    $\tilde{V} \leftarrow \{\text{LANESHIFT}(V_k, \delta)\}$ 
23:   for  $q \leftarrow k$  to  $G - 1$ ;  $\delta \in \text{ACTIVESHIFTS}(q, k)$  do
24:      $A \leftarrow \text{CLIENT.EMITWEIGHTTILE}(q, k, \delta)$ 
25:      $C_q \leftarrow C_q + A \odot \tilde{V}_\delta$ 
26:   end for
27: end for

```

▷ re-encrypt only

Stages 4 and 5: projection, MLP, residuals

```

28: for  $g \leftarrow 0$  to  $G - 1$  do
29:    $R_g \leftarrow c_g + \text{MATMUL}(\text{BABYROTATIONS}(C_g), W_{\text{proj}})$ 
30: end for
31: FLUSHENCODEDWEIGHTS
32: for  $g \leftarrow 0$  to  $G - 1$  do
33:    $\beta \leftarrow \text{BABYROTATIONS}(\text{LAYERNORM}(R_g, \gamma_2))$ 
34:    $h \leftarrow \sum_{c < 4} \text{MATMUL}(\beta, W_{\text{fc}}^c) \odot \text{copy}_c$ 
35:    $a \leftarrow \text{CLIENT.GELU}(h)$ 
36:    $c'_g \leftarrow R_g + \sum_{c < 4} \text{MATMUL}(\text{BABYROTATIONS}(a \odot \text{copy}_c), W_{\text{mlp}}^c)$ 
37: end for
38: return  $c'_0, \dots, c'_{G-1}$ 

```

Algorithm 2 The client boundary. The data owner is the only party holding sk , and every value it decrypts derives from its own query.

```

1: function CLIENTBOUNDARY( $c, f$ )
2:    $v \leftarrow \text{DECODE}(\text{DECRYPT}(\text{sk}, c))$ 
3:    $v' \leftarrow f(v)$ 
4:   return  $\text{ENCRYPT}(\text{pk}, \text{ENCODE}(v'))$ 
5: end function

```

▷ evaluated exactly, in double precision
▷ fresh ciphertext at level 0

Algorithm 3 LayerNorm. The server computes both statistics under encryption; only the inverse square root crosses the boundary. The released checkpoint carries no bias term.

```

1: function LAYERNORM( $x, \gamma$ )
2:    $\mu \leftarrow \text{REDUCESUM}(x) \odot (1/D)$ 
3:    $z \leftarrow x - \mu$ 
4:    $\sigma^2 \leftarrow \text{REDUCESUM}(z \odot z) \odot (1/D) + \varepsilon$ 
5:    $r \leftarrow \text{CLIENTBOUNDARY}(\sigma^2, t \mapsto t^{-1/2})$ 
6:   return  $(z \odot r) \odot \gamma$ 
7: end function

```

Algorithm 4 Packed matrix product, baby-step giant-step over the packed diagonals. $N_1 = N_2 = 32$ and $N_1 N_2 = W = 1024$.


```

1: function MATMUL( $\beta, M$ )
2:    $d \leftarrow \text{ENCODEDDIAGONALS}(M, \text{level}(\beta_0))$  ▷ encoded once, cloned per use
3:    $y \leftarrow 0$ 
4:   for  $j \leftarrow 0$  to  $N_2 - 1$  do
5:      $u \leftarrow \sum_{i < N_1} \beta_i \odot \text{CLONE}(d_{N_1 j + i})$ 
6:      $y \leftarrow y + (j = 0 ? u : \text{ROT}(u, B N_1 j))$ 
7:   end for
8:   return  $y$ 
9: end function

```


One ciphertext holds 32,768 slots

4 copies $\times$ 1,024 features $\times$ 8 token lanes = 32,768 slots. The copies carry the 4 $\times$ MLP expansion.



GSR: 600 bp become 103 tokens in 13 groups

600 bp $\div$ 6-mer = 100 tokens, plus 3 task specials = 103. 103 tokens across 8 lanes rounds up to 13 groups; the final group is partly padded.



The 13 groups give 91 lower-triangular causal score tiles, and just 1 of the 104 lanes is padding.

Figure 3: Slot layout for the genomic-signal task. Each ciphertext carries four copies of 1,024 features across eight token lanes: 768 features are active and the remainder supplies power-of-two padding and attention staging. The 103-token prompt occupies 13 groups and generates 91 lower-triangular causal score tiles.

The algorithms also expose the protocol transcript. Its order and size are determined by the circuit and sequence length, never by decrypted content. The evaluated block makes 857 physical client boundary crossings, including 91 score-tile decryptions and 727 weight-tile encryptions for attention. Attention therefore accounts for 818 of the 857 client calls, or 95.4%. In directional object accounting it contributes 818 of the block's 896 boundary ciphertext objects, or 91.3%. These homogeneous counts identify attention as the main boundary-schedule target without treating the implementation's mixed-unit logical-instance counter as scalar work. The implementation asserts the complete operation schedule, boundary counts, and remaining depth on every run and aborts when any differs. Section 6 can thus treat an unchanged circuit as a checked invariant rather than as an assurance.

5.4 Packing

Each ciphertext has 32,768 usable slots, indexed by activation copy, padded feature, and token lane. The layout assigns these slots as 4 copies by 1,024 features by 8 token lanes, with the token lane as the innermost index. Rotations by whole-lane strides can therefore move features without mixing tokens, while smaller rotations align tokens within a group. The released hidden width of 768 occupies part of the padded feature region; the remainder supplies the power-of-two width required by the linear transform and staging space for per-head attention scores. In dense transforms, rows and columns outside the active width contribute zeros; padding therefore affects cost without changing the numerical result.

The four copies match the MLP expansion from width 768 to 3,072. They let its expansion chunks occupy parallel regions of one ciphertext rather than requiring a separate wide layout. The 103 tokens fill 13 ciphertext groups, with padding in the final group. Within a dense transform, a 32×32 baby-step giant-step decomposition covers the 1,024 padded features; baby rotations are shared where the schedule permits. Packing eight tokens together reduces the number of dense encrypted products from 1,236 serial-equivalent products to 156, a structural reduction of 7.92 $\times$. No latency factor follows from that operation count alone.

5.5 Causal attention

Attention follows a chunked lower-triangular schedule over the 13 token groups. Only causal group pairs are visited, producing 91 encrypted score tiles. Shifts needed by several later query groups are materialized per key group and reused. The client collects all score tiles before applying stable softmax over complete causal rows; normalization is therefore a global synchronization point in this design. It then emits the encrypted weight tiles consumed by the context computation.

Plaintext model weights do not eliminate ciphertext–ciphertext multiplication inside attention: both query–key scoring and softmax-weight–value aggregation multiply one activation-derived quantity by another. Of the block’s 1,506 ciphertext–ciphertext multiplications, LayerNorm contributes 26 centered-value squares and 26 products with client-returned inverse square roots. Attention contributes 727 query–key products and 727 weight–value products, accounting for 97% of the total. The latter term follows directly from returning exact attention weights as fresh ciphertexts; it is the arithmetic price paid to keep the attention nonlinearity exact at the client boundary.

5.6 Depth and parameters

The retained CKKS context uses ring dimension 65,536, yielding 32,768 usable slots at a 128-bit classical security level. It uses 50-bit scaling, a 60-bit first modulus, 3 hybrid key-switching digits, and 80 rotation keys. Multiplicative depth 13 is the smallest demonstrated passing depth for this scale, digit count, ring, and packing. The block completes with 6 of the 13 levels remaining. The retained setting is the demonstrated feasible point in this parameter space; parameter optimization is a separate experiment.

5.7 Composition across blocks

Composition uses a declared full-hidden-state client refresh between blocks. The data owner decrypts the packed block output, restores its token-vector layout, and encrypts fresh level-0 inputs for the next block under the same context and key lineage. Two consecutive released blocks compose by this mechanism at two tokens with relative error 8.48×10^{-11} . Device-wide GPU memory used peaks in the first block and does not grow in the second.

The complete evaluation uses the same key lineage with 11 full-hidden-state refreshes and one last-token refresh into the task head; Section 7 reports the measured outcome.

361 6 Host-side systems work at fixed circuit

362 Encoding each distinct weight diagonal once removes redundant host work but creates a memory lifetime problem. An
363 implementation that retained every encoded template was terminated by the operating system at 61,863 MiB resident,
364 when the MLP began building templates on top of those cached by earlier stages. Bounding the cache to one live
365 stage makes reuse executable within the tested host-memory envelope. Peak host memory is then determined by the
366 largest stage, not by every encoded weight in the block. The remaining changes govern when host data are prepared,
367 how long they remain resident, and where CPU work executes.

368 Every change in this section holds the encrypted circuit fixed. The retained implementation executes the same 156
369 dense products, 177,734 ciphertext–plaintext multiplications, 1,506 ciphertext–ciphertext multiplications, 8,173 rota-
370 tions, and 857 client boundary crossings before and after the change. Packing, multiplicative depth, boundary order,
371 and the frozen plaintext reference are also unchanged. These properties are enforced at run time: the implementation
372 compares the observed schedule and boundary counts with their expected values, checks the remaining depth before
373 the MLP, and aborts on any deviation.

374 None of the retained changes removes a dense transform, reduces the number of attention tiles, moves linear algebra
375 to the client, or relaxes the acceptance criterion. The section consequently reports verified changes in executed work
376 and memory, but no elapsed-time comparison. No dedicated-node measurement of the pre-optimization configuration
377 exists, so the complete-model timings in Section 7 describe only the retained implementation.

378 6.1 Host-side encoding redundancy

379 The baseline rebuilt and encoded every packed diagonal for every dense product. Its 156 products each required 1,024
380 diagonal plaintexts, producing 159,744 encoding operations per block. About 90% of this work was byte-identical: the
381 block has 12 weight matrices, and each is used across all 13 token groups. The GPU could not begin the corresponding
382 multiplication until the host had prepared its plaintexts.

383 In the baby-step giant-step transform, a plaintext diagonal is not merely a view of the original matrix. It must be
384 packed, scaled, and encoded for the ciphertext level at which multiplication will occur. Repeating that conversion
385 for every token group made a model-invariant representation step dominate the call path. Caching therefore targets
386 plaintext representation while preserving the encrypted products and accumulations.

387 The operation counts locate the redundancy without relying on elapsed time: identical weight diagonals are rebuilt
388 across token groups even though the public matrix does not change. The optimization target is the lifetime and
389 placement of plaintext preparation while encrypted matrix arithmetic remains fixed.

390 6.2 Encode once, clone per use

391 The retained cache encodes each distinct weight diagonal into a pristine host template. A multiplication receives a
392 lightweight clone that shares the encoded host data by reference while carrying fresh device state; the clone is loaded
393 and evicted for that operation. This removes 90% of the 159,744 encoding operations without keeping a plaintext
394 object resident on the GPU between multiplications.

395 Cloning is required by the evaluated backend’s object semantics. Its plaintext-load path returns early when an object
396 is already marked as loaded, so directly reusing a template across multiplicative levels can pair a new ciphertext level
397 with stale residue-number-system limbs. Fresh device state on each clone forces a load at the level required by the
398 current operation. An isolated GPU micro-test checked clone-versus-fresh-encode equality, cross-level reuse, stable
399 device memory under repeated multiplication, and equality between parallel and serial encoding before the cache was
400 applied to the complete block.

401 Templates are indexed by both weight matrix and multiplicative level. This distinction preserves the backend’s level
402 contract while still sharing work among token groups that reach the same matrix at the same level. The cache therefore
403 stores reusable encoded data, not a ciphertext or a mutable GPU object, and leaves homomorphic multiplication itself
404 unchanged.

405 6.3 Bounding host memory

406 Removing repeated encoding exchanges computation for host memory. An early cache retained templates from every
407 completed stage, and the process did not survive the MLP. Two exact changes bound the cache. First, a template
408 is encoded directly at its multiplicative level of use, so it carries only the limbs needed by the receiving ciphertext.
409 This changes no arithmetic because the backend would otherwise drop a level-0 plaintext to the same level before
410 multiplication. Second, the implementation flushes the cache after query/key/value projection and after the attention

Bounded-cache run: stage-local host-memory peaks




Figure 4: The bars are stage RSS measurements from one 103-token block with bounded caching. Flushing limits the live cache to the largest stage. The red line records a separate unbounded-cache attempt that the operating system killed at 61,863 MiB.

projection. No weight matrix recurs across either boundary, so retaining those templates cannot produce another hit. The flush points are the final use sites for their respective weight sets.

The lifetime policy follows the block topology. Within a stage, the group-major loop revisits its weights as it advances through token groups, so early eviction would discard real reuse. Across stages, the next operation consumes a different set of weights; eviction at that point releases memory without creating a later re-encoding. The cache consequently remains large enough to serve the schedule and no larger.

The target-process resident-set trace follows this policy. Query/key/value projection reaches 25,530 MiB. Attention begins after a flush and uses 12,434 MiB. The MLP holds eight matrices and reaches the block peak of 48,610 MiB. The largest live stage, rather than every encoded weight in the block, determines peak host memory.

6.4 Placement and parallelism

Plaintext preparation also depends on where host work runs. The retained implementation binds threads to physical cores and uses NUMA-local allocation so encoding threads stay close to the memory they allocate and to buffers staged for the GPU. Decryption, exact nonlinear evaluation, and re-encryption use the same host resources. Their elapsed-time effect awaits a sourced paired measurement.

Template construction was then batched across the available 32 cores with dynamic scheduling. Each diagonal can be encoded independently, making this step parallel without changing the order or result of encrypted multiplications. Parallel and serial template construction were checked for equal outputs before the parallel path was retained.

Affinity and batching address different parts of the same host path: affinity controls placement, while batching exposes independent diagonals to the available cores. Together they make plaintext preparation a core-local parallel batch under the same cryptographic parameters.

6.5 Remaining optimization targets

Mask plaintexts are still encoded afresh at approximately 17,400 uses per block. Head masks, score-isolation masks, lane masks, and copy-activity masks all have fixed values under the declared layout, so host templates could remove the same class of repeated work already eliminated for weight diagonals. This target remains unimplemented.

A second opportunity lies in the four packed activation copies. They currently carry duplicates through a dense transform. An analyzed layout would apply distinct transforms to those copies, placing query, key, and value projections together and likewise combining the MLP expansion chunks. At fixed model arithmetic, the projected schedule

Table 4: Host-side changes at fixed circuit. Every row leaves the operation schedule, packing, multiplicative depth, and frozen plaintext reference unchanged. Effects are verified changes in executed work or in memory; the table carries no elapsed-time column, because no dedicated-node measurement of the pre-optimization configuration exists.

Change	Verified effect	Status
Re-encode every plaintext per multiplication	159,744 encodings	baseline
Thread affinity, NUMA-local allocation	host placement changed	retained
Encode once, clone per use	removes 90% of encodings	retained
Parallel batched encoding	equal to serial encoding	retained
Encode at use level, flush between stages	peak 48,610 MiB RSS	enabling
Net	checked operation schedule unchanged	

reduces dense products from 156 to approximately 52 and dense ciphertext–plaintext multiplications from 159,744 to approximately 53,248; the total ciphertext–plaintext count falls by roughly 60% before mask and copy-repair overhead. This is a derived projection for an unimplemented layout, and it requires an exact reconstruction proof before implementation.

Client work has a separate opportunity. Attention accounts for 818 of the 857 client calls and 818 of the 896 boundary ciphertext objects per block. Each attention tile’s cryptographic work is independent. Concurrent client boundaries were unsafe under the evaluated shared crypto context, so the retained implementation executes them sequentially. A thread-safe context would expose this independent work to concurrency.

7 Results

The client-assisted circuit is evaluated along four axes: numerical agreement, executed work, memory, and complete-inference cost.

7.1 Evaluation setup and endpoints

The encrypted input is selected deterministically as the first record in the verified positive genomic-signal FASTA input. It is one 600-base-pair fragment, not a random, margin-stratified, or task-set sample. Local tokenization and embedding produce the 103-token prompt before encryption; the measured encrypted endpoint contains all 12 released blocks and the released task head.

The complete model is executed once on a whole-node allocation with one 40 GB A100 GPU and 32 CPU threads; no contention is detected in the recorded telemetry, although scheduler exclusivity was not requested. Timing starts before weight and reference loading and ends after final decryption. Server and client roles share one process and node, so boundary calls include cryptographic work but no serialization or network transport. That single execution supplies every complete-model timing below; it supports neither a mean nor a service rate. Earlier contended timings are excluded.

7.2 Correctness

Correctness follows the three-link chain in Section 4. The released float32 path is recorded separately. The independent NumPy float64 reference is then checked against a float64-cast PyTorch shadow at 2×10^{-5} for all 12 blocks and the head, after which encrypted execution is compared with the NumPy reference. This distinguishes model-dtype drift from CKKS error and prevents an unchecked reimplementaion from judging its encrypted counterpart.

For one block at the 103-token task prompt, the decrypted output has relative infinity-norm error 4.64×10^{-9} against an acceptance tolerance of 4×10^{-2} fixed before the run, leaving seven orders of magnitude of headroom. The worst token-relative error is 4.7×10^{-9} , and the largest absolute deviation is 7.44×10^{-8} . These values establish numerical agreement for the encrypted block with the released weights and exact client nonlinearities.

The complete model is held to the same unchanged tolerance. Evaluating all 12 released blocks and the released task head at the 103-token task prompt yields a decrypted classification margin that differs from the manual NumPy float64 margin by 8.56×10^{-9} in relative terms and 1.05×10^{-7} in absolute terms. The encrypted, manual float64, and released float32 paths all produce label N for this prompt; relative to the released float32 margin, the encrypted margin differs by 1.61×10^{-6} .

Error accumulates with depth. The hidden-state relative infinity norm is 4.09×10^{-9} after the first block and 2.07×10^{-8} after the twelfth, while the largest absolute deviation grows from 6.56×10^{-8} to 7.97×10^{-6} , roughly two orders of

Table 5: Checked operation schedule for the evaluated transformer block.

Operation	Count	Executor
Dense products	156	compute provider
Ciphertext–plaintext multiplications	177,734	compute provider
Ciphertext–ciphertext multiplications	1,506	compute provider
Rotations	8,173	compute provider
Slot reductions	8,776	compute provider
Physical client boundary crossings	857	data owner
Attention-related client calls	818 (95.4%)	data owner

Table 6: Derived ciphertext-object accounting for a complete inference. Counts follow from the checked protocol schedule; they are not observed network messages.

Direction or quantity	Count	Scope
Server to client	1,719	nonlinear and refresh boundaries
Client to server	9,339	nonlinear and refresh boundaries
Both directions	11,058	boundary ciphertext objects
Initial encrypted inputs	13	client to server
Final encrypted output	1	server to client
Total including endpoints	11,072	ciphertext objects
Minimum sequential phases	63	boundaries, with ideal same-stage batching
Minimum phases including endpoints	65	algorithmic lower bound
Serialized byte volume	unknown	no wire format was measured

magnitude across the model. Growth of that rate leaves the twelfth-block activation six orders of magnitude inside the acceptance tolerance. The classification-margin error is 6.67 orders below the same threshold. Approximation error accumulates, but it does not change the decision for the evaluated prompt.

7.3 Executed work

The implementation asserts every count in Table 5, the boundary sequence, and the remaining multiplicative depth before the MLP. Any deviation aborts the run. This turns the operation schedule into evidence that the optimized implementation executes the same circuit, not an assurance inferred from a passing output.

Attention dominates the homogeneous boundary-schedule measures: it accounts for 818 of the 857 client calls per block and 818 of the 896 ciphertext objects sent across block boundaries. Boundary optimization is therefore principally an attention problem, even though the protocol contains three nonlinearity classes. The implementation also checks a heterogeneous logical-instance counter, but its LayerNorm, GELU, and attention increments have different units and are not interpreted as scalar work or traffic.

The dense schedule also records the benefit of token packing. Eight-token packing reduces the serial-equivalent dense-product count from 1,236 to 156, a $7.92\times$ reduction in operations. Rotations, masking, client work, and host preparation remain in the packed circuit, so the factor describes operation count rather than latency.

One block presents 130 ciphertext objects to the client and receives 766 in return. Across the complete model, ideal batching of independent calls within a protocol stage yields a lower bound of 65 sequential transport phases including input and output. The present implementation instead performs 10,299 in-process boundary calls. Ciphertext levels, serialization representation, compression, and transport overhead were not recorded, so a byte count cannot be derived honestly from the object count.

7.4 Memory

Device-wide GPU memory used during the evaluating run reaches 9839 MiB, within the 40 GB device. The GPU was empty at allocation, making attribution to the run reasonable, but the sampler does not expose a per-process memory metric.

Host memory follows the lifetime of encoded plaintext templates. Query/key/value projection reaches 25,530 MiB of target-process resident memory. A cache flush precedes attention, whose recorded resident memory is 12,434 MiB. The MLP holds the largest live weight set and reaches the block peak of 48,610 MiB. These stage transitions match the explicit flush points in Section 6; encoded weights do not accumulate across the entire block.

Sampled GPU use stays below saturation; memory remains bounded across 12 blocks
The plotted curve is thinned from the full 653-sample trace.




Figure 5: Resource trace of the complete encrypted inference. Sampled GPU utilization reaches a maximum of 51%, showing underutilization without assigning its cause. Host memory repeats the stage pattern of Section 6; device-wide GPU memory used reaches its peak during the first block and does not grow with block index. The plotted curve is thinned from the full 653-sample trace; the twelve cycles are the twelve blocks.

The two memory results have different mechanisms. The non-interactive composition configuration exceeded its tested 80 GB GPU envelope while loading the deep context and evaluation keys. Client re-encryption shortens the encrypted segments and resets their depth budget, moving that accelerator-memory barrier. Stage-bounded caching prevents the resulting host-side plaintext templates from becoming an unbounded second barrier.

The complete model confirms both bounds at full depth. Device-wide GPU memory used rises during context and key loading, reaches 9839 MiB after roughly 400 s, and then holds that value exactly for the remainder of the run. This matches the single-block footprint, with no growth by block index. Host resident memory peaks at 50,930 MiB, consistent with the MLP-stage maximum recurring once per block instead of accumulating. Client assistance therefore removes the accelerator-memory barrier that stopped the evaluated non-interactive configuration, and the flat device profile across twelve blocks shows this follows from the depth bound rather than from the particular input.

7.5 Composition and evaluation boundary

Two consecutive released blocks compose at reduced sequence length through a full-hidden-state client refresh, reaching 8.48×10^{-11} relative error after the second block. Device-wide GPU memory used peaks in the first block and does not increase in the second. This experiment checks the mechanism used for deeper composition: a fresh encryption starts the next block at level 0 without carrying forward the preceding block’s multiplicative depth or GPU-resident temporaries.

The complete model uses that mechanism throughout. All 12 blocks and the task head run in one key lineage, separated by 11 full-hidden-state refreshes between consecutive blocks and one further refresh of the final token into the head. Multiplicative depth stays at 13 throughout, the same budget the single block requires, and the run performs no homomorphic bootstrap at any point. Depth is thus a property of the longest segment between two client boundaries rather than of network depth, which is what makes the twelve-block key lineage expressible at all under these parameters.

Table 7: Encrypted evaluation completed at each reported scope. Hidden-state error is a relative infinity norm over the activation; margin error is the scalar the classifier acts on. The two are different quantities and are kept in separate columns.

Scope	Status	Hidden-state error	Margin error
One block, 2 tokens	measured	—	—
Two consecutive blocks, 2 tokens, client refresh	measured	8.48×10^{-11}	—
One block, 103 tokens (task length)	measured	4.64×10^{-9}	—
Twelve blocks and task head, 103 tokens	measured	2.07×10^{-8}	8.56×10^{-9}
Encrypted token-index lookup	out of scope	—	—

Encrypted evaluation: 6,595 s of 6,683 s wall clock

One measured 12-block + task-head execution at the 103-token GSR prompt.




Figure 6: Encrypted-evaluation time for one complete inference. The provider-side bucket is the residual after client time and includes GPU-backend work plus host preparation, synchronization, and evaluator overhead. The data-owner bucket includes declared nonlinear boundaries and full-hidden-state refreshes on CPU with no GPU. The two shares sum exactly to encrypted evaluation; setup and finalization remain in Table 8.

7.6 Cost of a complete encrypted inference

A complete encrypted inference takes 6683 s of wall clock, or 1.86 hours. Encrypted evaluation accounts for 6594.748 s of that, and setup is negligible at this scale: context construction and key generation together take under four seconds, and the single final decryption takes 0.044 s. The instrumented phases are a decomposition rather than a partition, leaving 59.504 s, or 0.9% of wall clock, outside them; Table 8 shows that remainder rather than distributing it silently.

Within encrypted evaluation the split between the two parties is exact: the residual provider-side bucket uses 5638.561 s and the data owner’s boundary work uses 956.187 s, and the two sum to the evaluation total without residue. The provider bucket includes GPU-backend work and the host support around it; it is not GPU-kernel time. The compute provider therefore carries 85.5% of the encrypted-evaluation interval and the data owner 14.5%, with the data owner’s share running on CPU alone. Across 10,299 in-process boundary calls the mean call costs 0.093 s.

The twelve blocks account for 6565.993 s and the encrypted task head for 9.046 s, so the head is not a meaningful part of the cost. The twelve refreshes together cost 19.710 s, which places the price of the composition mechanism at roughly 0.3% of encrypted evaluation. Dividing block evaluation by twelve gives 547.166 s per block. The twelve individual block times span 533.869 s to 558.943 s, a spread of 4.7% across architecturally identical blocks on one node. That bounds within-run scheduling jitter; it is not run-to-run variance, which a single execution cannot report.

Table 8: One complete encrypted inference: twelve blocks and the released task head at the 103-token task prompt, on one A100 GPU node using 32 CPU threads. Single execution on a whole-node allocation with no contention detected in the recorded telemetry; these are one sample’s values, not averages over repetitions. Shares are rounded independently and need not sum to 100%.

Phase	Seconds	Share of wall
Load released weights and verified reference	23.494	0.4%
Crypto context, key generation, GPU key upload	3.945	0.1%
Encrypt the embedded input	1.265	0.0%
Encrypted evaluation	6594.748	98.7%
provider-side encrypted evaluation	5638.561	84.4%
declared client boundaries and refreshes (CPU)	956.187	14.3%
Final decrypt	0.044	0.0%
Uninstrumented remainder	59.504	0.9%
Wall clock	6683	




Figure 7: Causal attention score-tile schedules across the encrypted-scope tasks. Counts follow from ciphertext groups rather than from a timing model; only genomic-signal recognition has a measured complete inference.

7.7 Sequence-length structure

Sequence length follows from each task format. Genomic-signal recognition uses 600 base pairs, which become 100 sequence tokens under 6-mer tokenization; three task-format tokens produce the 103-token encrypted anchor. Eight-token packing maps that prompt to 13 ciphertext groups. The shorter promoter and splice-site prompts use fewer groups under the same layout, yielding smaller causal schedules.

The abundance-regression task lies outside the encrypted scope. Its 10,500 base pairs produce 1,755 tokens and 220 ciphertext groups under the same tokenization rule. The causal attention schedule grows quadratically in the number of groups, placing this task in a different circuit-size regime from the classification prompts.

7.8 Feasibility

Complete-model graph feasibility is established under the stated protocol. All 12 released blocks and the task head execute at the task’s 103-token prompt under one key lineage with declared refreshes. Multiplicative depth remains 13, device-wide GPU memory used remains 9839 MiB after loading, and no homomorphic bootstrap is performed. These are measurements of the complete computation, not projections from a block.

Numerical generality has a narrower scope. The one selected prompt reproduces the frozen label with 8.56×10^{-9}

relative margin error, but one prompt does not establish encrypted task accuracy or input-independent stability. The affirmative result is that the executed case meets the predeclared numerical criterion while composing the complete released graph.

Encrypted token-index lookup and network transport remain outside this verdict, as detailed in Section 9.

7.9 Practicality, judged separately

Feasibility does not make the current implementation practical. One encrypted classification of one 600-base-pair genomic fragment occupies an A100 node for 1.86 hours and requires an online key holder throughout, which rules out interactive per-query use in the evaluated form. Suitability for a specified non-interactive workload has not been measured.

Cross-paper results provide context rather than a controlled verdict. THOR reports 602 s for a non-interactive encrypted forward pass of a 12-layer, 768-wide encoder at 128 tokens on a single A100, at the ring degree, multiplicative depth, and security level used here, with no client participation at any point Moon et al. [2025]; the 6683 s measured here is roughly an order of magnitude above it at a shorter prompt. Model semantics, nonlinear treatment, packing, software, and measurement endpoints differ, so this implementation is not performance-leading and the gap does not isolate the cost of client assistance; Section 8 sets out the differences.

The measured trace does, however, indicate where that cost sits. Sampled GPU utilization has a median of 32% and a maximum of 51%, so the accelerator is not the saturated resource, even though the trace does not attribute wall time among encoding, synchronization, memory behavior, kernel efficiency, and boundary work. Read with the 85.5/14.5 provider–client split and the 0.3% cost of the refresh mechanism, this places the optimization target in the provider-side encrypted path rather than in the interaction the protocol adds. The practicality verdict is nonetheless negative for the evaluated form: it is too slow for interactive use, and throughput, serialization, and networked deployment remain unmeasured.

8 Related work

8.1 Private transformer inference

Private neural-network inference has long combined cryptographic tools according to the operation being evaluated. Gazelle established an influential hybrid design in which homomorphic encryption handles linear layers and garbled circuits handle nonlinearities Juvekar et al. [2018]. Transformer systems inherited this decomposition but had to contend with attention, normalization, and activation functions at substantially greater scale. BOLT, Iron, Nimbus, and BumbleBee combine homomorphic encryption with secret sharing, oblivious transfer, and specialized two-party protocols for transformer inference Pang et al. [2024], Hao et al. [2022], Li et al. [2024], Lu et al. [2025]. Their protocols optimize a joint secure computation between parties; the present work instead keeps the server’s linear path in CKKS and assigns exact nonlinear evaluation to the data owner at explicit boundaries. This choice changes both the security contract and the allocation of work between the parties.

The client-aided division of labor predates transformer work. nGraph-HE2 evaluates standard pre-trained networks under CKKS while returning to the client to compute their native activation functions, precisely so that those activations need not be replaced by polynomials Boemer et al. [2019]. THE-X carries the pattern to transformers: the server evaluates HE-compatible operations and the client decrypts and re-encrypts around nonlinear functions Chen et al. [2022]. Safire is the closest architectural antecedent Biswas et al. [2025]. It likewise places homomorphically evaluated linear layers on the server, then has the client decrypt an intermediate, apply the nonlinearity in plaintext, and return a fresh encryption. Safire studies convolutional networks and adds randomized output permutation to make reconstruction from released intermediates harder. The client-assisted pattern is thus established prior art, and we adopt rather than propose it. What this study adds is its application to a complete forward pass of a released genomic transformer at that model’s task-defined prompt, preserving the released weights and exact nonlinearities, together with an independently checked plaintext reference that makes the resulting arithmetic, boundary schedule, and memory behavior reproducible.

The non-interactive line of work supplies useful cost context. THOR evaluates a 12-layer, 768-wide encoder at 128 tokens on a single A100, at ring degree 2^{16} , depth 13, and 128-bit security, in 602 s end to end, with no client interaction at any point Moon et al. [2025]. Its headline configuration overlaps with this study in several dimensions, but the workload, protocol, and implementation differ from the 6683 s execution measured here. NEXUS reaches non-interactivity over RNS-CKKS with SIMD ciphertext compression and a secure argmax Zhang et al. [2025]; its headline per-inference figure is amortized over a batch, and its single-prediction latency is substantially higher, so the two readings should not be conflated. MOAI reports a further 52.8% reduction in evaluation time against THOR by

607 holding one packing convention across all layers and removing rotations from softmax and normalization [Zhang et al.](#)
608 [\[2026\]](#).

609 These are uncontrolled cross-system comparisons. Model semantics, prompt length, nonlinear circuits, bootstrapping,
610 packing, libraries, CPU paths, and measurement endpoints differ. They establish that the present implementation is
611 not performance-leading, but they do not isolate the cost of client assistance. The evaluated design occupies a different
612 point in that space: it retains the released nonlinear formulas exactly, requires no homomorphic bootstrap, and pairs
613 its execution with a frozen, independently checked plaintext reference.

614 Powerformer represents a complementary model-modifying direction, training HE-suitable polynomial components
615 rather than preserving the released transformer unchanged [Park et al. \[2025\]](#). ELLMo further shows how packing and
616 depth-aware design can restructure encrypted transformer inference [Guzelhan et al. \[2026\]](#). These systems reinforce
617 that performance comparisons must account for model and circuit changes, not only headline dimensions.

618 8.2 Homomorphic encryption for genomics

619 Homomorphic encryption already has a substantial genomic-computing lineage. The iDASH secure genome analysis
620 competition evaluated encrypted association studies, sequence search, and related privacy-preserving workflows [Kuo](#)
621 [et al. \[2020\]](#). Private genomic-query systems combine homomorphic encryption with hashing and set-intersection
622 techniques to test for variants without disclosing either the stored genome or the query [Çetin et al. \[2017\]](#). Those
623 workloads center on searches or prescribed statistics over genomic records. DNAGPT inference instead composes
624 dense projections, attention, normalization, and a feed-forward network while preserving the numerical behavior of a
625 released foundation model. This study extends the genomic HE lineage to the feasibility boundary of that transformer
626 forward pass.

627 The iDASH series has since moved from encrypted statistics toward encrypted neural evaluation, and supplies the
628 closest prior art to this study. Its 2024 homomorphic-encryption track distributed DASHformer, a protein-family
629 classifier comprising a tokenizer, an encoder, a single transformer block with four-head self-attention, and a dense
630 classification layer, totaling 138,905 parameters over 50-residue fragments [Harmanci et al. \[2024\]](#). The winning solution
631 evaluates it non-interactively under leveled CKKS, with no bootstrapping and no client participation at any point,
632 classifying a batch of 163 encrypted sequences in about 165 s at 128-bit security [Yang et al. \[2024\]](#). On interaction
633 that result is strictly stronger than the design evaluated here, which returns to the key holder at every nonlinearity. It
634 reaches that strength under different constraints: one transformer block against twelve, a purpose-trained competition
635 classifier against a released 0.1-billion-parameter backbone, and nonlinearities replaced by data-driven non-polynomial
636 fits rather than preserved. To our knowledge, no released, pretrained genomic foundation model has previously been
637 evaluated under homomorphic encryption for a complete forward pass.

638 8.3 GPU-accelerated homomorphic encryption

639 GPU CKKS libraries move number-theoretic transforms, key switching, rotations, and ciphertext arithmetic onto
640 accelerators; the evaluated backend supplies those primitives and interoperates with the host cryptographic run-
641 time [Agulló-Domingo et al. \[2025\]](#), [Al Badawi et al. \[2022\]](#). Its API also shaped the available systems designs. In
642 particular, the absence of ciphertext serialization prevented process-sharded evaluation, leaving thread-level concur-
643 rency and in-process device placement as the implementable paths for this implementation.

644 Recent multi-GPU transformer work treats placement and communication-computation overlap as first-class schedul-
645 ing problems [Gong et al. \[2026\]](#). That literature motivates the placement and overlap options discussed in Section 6.
646 The current implementation neither serializes ciphertexts between processes nor exercises a networked client/server
647 transport, so those paths must be evaluated directly before adopting a multi-device schedule.

648 8.4 Homomorphic matrix primitives

649 Matrix layout determines which rotations, plaintext encodings, and reductions an encrypted linear layer requires.
650 THOR’s diagonal-major representation and compact packing specialize that layout for transformer inference [Moon](#)
651 [et al. \[2025\]](#); HE-BLAS develops a broader library of homomorphic linear-algebra kernels and shape-aware algo-
652 rithms [Bae et al. \[2026\]](#). These works motivate the remaining targets in Section 6, as does MOAI’s rotation-free
653 formulation of softmax and normalization [Zhang et al. \[2026\]](#). Those targets include reducing repeated packing work,
654 copying fewer distinct diagonal plaintexts, and replacing segmented slot reductions with kernels matched to the tensor
655 shapes. ELLMo likewise treats packing and depth as coupled optimization choices [Guzelhan et al. \[2026\]](#). Because
656 packing, key material, level consumption, and the operation schedule interact, each candidate primitive must be
657 revalidated against the frozen plaintext reference and checked schedule before its performance is compared.

9 Limitations

The complete model has been evaluated once. All 12 released blocks and the task head run at the task’s full 103-token prompt on a whole-node allocation with no contention detected in the recorded telemetry, and the reported timings come from that single execution. One sample supports a statement about the cost of that inference; it does not support a mean, a variance, or a service-rate claim, and the run has not been repeated. Nor does the study report encrypted task accuracy: the encrypted prediction is checked against the frozen plaintext reference for one input, not evaluated over a test set. That input was selected deterministically as the first record of the positive genomic-signal input, not randomly or by a predeclared margin stratum.

Interaction is a standing requirement rather than an implementation detail, and it is the limitation that the cost comparison in Section 8 makes concrete. The protocol needs an available key holder at every one of the 10,299 declared boundaries of a complete inference; a data owner that goes offline mid-evaluation has no way to finish it, and no party other than the data owner can complete a boundary. Non-interactive designs carry no such dependency. This study therefore establishes no advantage over them in latency, throughput, or availability, and the cross-system comparison in Section 8 is not controlled.

Client and server are roles inside one program on one node. Boundary calls perform the cryptographic operations directly, without ciphertext serialization, network transport, or a key custody service. A physical client boundary crossing is therefore not a network message. A derived schedule accounts for ciphertext objects and minimum dependency phases, but serialized byte volume, bandwidth sensitivity, transport latency, and production key handling remain unknown.

Encryption begins at embedded numeric vectors. Tokenization, token-index lookup in the embedding table, and construction of the initial embedding occur before the encrypted boundary. The evaluated deployment assumes that the tokenizer and embedding table are distributed to the client. Private lookup is required if a provider will not make those components available.

Model confidentiality remains operational rather than cryptographic. Because the data owner sees full activations at every nonlinearity, the transcript divides the network into shallow segments with known nonlinearities. This creates a chosen-query extraction surface, and some affine segments may be identifiable from sufficiently diverse queries. Attention exposes composite relationships, however, and this study provides neither an extraction experiment nor a general recovery claim. Server-side weight placement may discourage casual copying, but it is not a cryptographic model-protection mechanism.

CKKS decryption is approximate. The client never returns a decoded value to the server, only a freshly encrypted result, but the server selects the ciphertext supplied to each declared boundary. Security analyses of approximate homomorphic encryption show why adversarially selected decryption queries require care [Li and Micciancio \[2021\]](#). The semi-honest threat model rules out deviations from the declared schedule, and the current re-encryption path applies no noise flooding. The measured single-block error of 4.64×10^{-9} leaves seven orders of magnitude inside the 4×10^{-2} acceptance tolerance, and the twelfth-block hidden state leaves six, but numerical headroom alone does not establish a secure flooding parameter. Selecting and validating that budget remains future work.

The adversary exclusions are defined in Section 3. The results do not extend beyond that threat model.

10 Conclusion

This study gives a split answer to the question of encrypted DNAGPT inference. Client-assisted CKKS executes all twelve released blocks and the task head at the 103-token prompt from encrypted embedded vectors. For the one input evaluated, the decrypted label matches the frozen reference and its margin has 8.56×10^{-9} relative error against the 4×10^{-2} predeclared tolerance. Graph feasibility is established for the complete released computation; numerical generality beyond this input is not. Client refresh bounds depth by the longest encrypted segment, device-wide GPU memory used stays at 9839 MiB across the blocks, and no homomorphic bootstrap is required. A stage-bounded plaintext cache separately prevents host-memory accumulation.

Cost is the constraint that remains. One encrypted classification of one 600-base-pair sequence was measured at 1.86 hours on an A100 GPU node, which forecloses interactive use in the evaluated form. It is a single-process, single-input, single-execution result rather than a service-rate or network measurement, and sampled GPU utilization peaking at 51% indicates unclaimed headroom without identifying its cause. The two verdicts are therefore separate, and both are informative: complete-model feasibility is established for the executed case, practicality is not, and the obstacle between them is engineering throughput on the provider’s encrypted path rather than correctness, multiplicative depth, or accelerator memory.

710 Future work

711 Further empirical work would require additional prompts, repeated timings, and a paired pre-optimization measure-
712 ment; none is inferred from the present run. Systems work should extend the encode-once strategy from weight
713 diagonals to mask plaintexts. Distinct transforms across the packed activation copies are the larger circuit oppor-
714 tunity. Before mask and copy-repair overhead, the derived schedule for that layout projects a reduction of roughly
715 60% in the block’s total ciphertext–plaintext count. A thread-safe cryptographic context would permit concurrent
716 independent client boundaries. Protocol closure also requires a noise-flooding budget for re-encryption. A deployment
717 study must add serialized network transport and private token-index lookup before measuring the complete service.

718 Ethics statement

719 This study involved no human participants, specimens, or private clinical records. It used public reference datasets
720 with recorded provenance. The work makes no clinical-performance claim and does not present the prototype as a
721 deployable privacy guarantee beyond its stated threat model.

722 Data availability statement

723 All evaluation datasets are public. Their original sources, versions, preprocessing, and any recovery route used are
724 recorded in the repository’s data-provenance document. Sequence data, model weights, and cryptographic keys are
725 not redistributed.

726 Code availability statement

727 The evaluation harnesses, encrypted-evaluation implementation, environment definitions, and the evidence record
728 for the paper are maintained at https://github.com/Georgakopoulos-Soares-lab/private_genomic_dnagpt. No
729 repository-level reuse license has yet been assigned.

730 References

- 731 V. Agarwal and J. Shendure. Predicting mrna abundance directly from genomic sequence using deep convolutional
732 neural networks. *Cell Reports*, 31(7):107663, 2020. doi: 10.1016/j.celrep.2020.107663.
- 733 C. Agulló-Domingo, O. Vera-López, S. Guzelhan, L. Daksha, A. El Jerari, K. Shivdikar, R. Agrawal, D. Kaeli,
734 A. Joshi, and J. L. Abellán. FIDESlib: A fully-fledged open-source fhe library for efficient ckks on gpus. In
735 *2025 IEEE International Symposium on Performance Analysis of Systems and Software*, pages 365–367, 2025. doi:
736 10.1109/ISPASS64960.2025.00045. URL <https://arxiv.org/abs/2507.04775>.
- 737 A. Al Badawi, J. Bates, F. Bergamaschi, D. B. Cousins, S. Erabelli, N. Genise, S. Halevi, H. Hunt, A. Kim, Y. Lee,
738 Z. Liu, D. Micciancio, I. Quah, Y. Polyakov, R. V. Saraswathy, K. Rohloff, J. Saylor, D. Suponitsky, M. Triplett,
739 V. Vaikuntanathan, and V. Zucca. OpenFHE: Open-source fully homomorphic encryption library. In *Proceedings*
740 *of the 10th Workshop on Encrypted Computing & Applied Homomorphic Cryptography*, pages 53–63, 2022. doi:
741 10.1145/3560827.3563379.
- 742 Y. Bae, J. H. Cheon, G. Hanrot, J. H. Park, and D. Stehlé. Fast homomorphic linear algebra with BLAS. *Journal of*
743 *Cryptology*, 39:25, 2026. doi: 10.1007/s00145-026-09580-x. Preprint: arXiv:2503.16080.
- 744 S. Biswas, P. Chartier, A. Dhasade, T. Jurien, D. Kerriou, A.-M. Kerrmarec, M. Lemou, F. Tranie, M. de Vos, and
745 M. Vujasinovic. Practical and private hybrid ml inference with fully homomorphic encryption. *arXiv preprint*
746 *arXiv:2509.01253*, 2025. URL <https://arxiv.org/abs/2509.01253>.
- 747 F. Boemer, A. Costache, R. Cammarota, and C. Wierzynski. nGraph-HE2: A high-throughput framework for neural
748 network inference on encrypted data. In *Proceedings of the 7th ACM Workshop on Encrypted Computing & Applied*
749 *Homomorphic Cryptography*, pages 45–56, 2019. doi: 10.1145/3338469.3358944. URL [https://eprint.iacr.org/](https://eprint.iacr.org/2019/947)
750 [2019/947](https://eprint.iacr.org/2019/947).
- 751 G. S. Çetin, H. Chen, K. Laine, K. Lauter, P. Rindal, and Y. Xia. Private queries on encrypted genomic data. *BMC*
752 *Medical Genomics*, 10(Suppl 2):45, 2017. doi: 10.1186/s12920-017-0276-z.

- T. Chen, H. Bao, S. Huang, L. Dong, B. Jiao, D. Jiang, H. Zhou, J. Li, and F. Wei. THE-X: Privacy-preserving transformer inference with homomorphic encryption. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 3510–3520, 2022. doi: 10.18653/v1/2022.findings-acl.277. URL <https://aclanthology.org/2022.findings-acl.277/>.
- J. H. Cheon, A. Kim, M. Kim, and Y. Song. Homomorphic encryption for arithmetic of approximate numbers. In *Advances in Cryptology – ASIACRYPT 2017*, volume 10624 of *Lecture Notes in Computer Science*, pages 409–437, 2017. doi: 10.1007/978-3-319-70694-8_15.
- Y. Erlich, T. Shor, I. Pe’er, and S. Carmi. Identity inference of genomic data using long-range familial searches. *Science*, 362(6415):690–694, 2018. doi: 10.1126/science.aau4832.
- European Parliament and Council of the European Union. Regulation (EU) 2016/679 of the european parliament and of the council. Official Journal of the European Union, 2016. URL <https://eur-lex.europa.eu/eli/reg/2016/679/oj>.
- Z. Gong, R. Ran, F. Yao, and W. Wen. AEGIS: Scaling long-sequence homomorphic encrypted transformer inference via hybrid parallelism on multi-GPU systems, 2026. URL <https://arxiv.org/abs/2604.03425>. Accepted at ICS 2026.
- S. N. Guzelhan, L. Daksha, C. Agulló Domingo, G. Jonatan, J. Kim, J. L. Abellan, D. Kaeli, and A. Joshi. ELLMo: Packing- and depth-aware encrypted transformer inference. IACR Cryptology ePrint Archive, Report 2026/198, 2026. URL <https://eprint.iacr.org/2026/198>.
- M. Gymrek, A. L. McGuire, D. Golan, E. Halperin, and Y. Erlich. Identifying personal genomes by surname inference. *Science*, 339(6117):321–324, 2013. doi: 10.1126/science.1229566.
- M. Hao, H. Li, H. Chen, P. Xing, G. Xu, and T. Zhang. Iron: Private inference on transformers. In *Advances in Neural Information Processing Systems*, volume 35, pages 15718–15731, 2022. doi: 10.52202/068431-1143.
- A. Harmanci, L. Chen, M. Kim, and X. Jiang. Descriptor: Benchmarking secure neural network evaluation methods for protein sequence classification (iDASH24). *IEEE Data Descriptions*, 1:109–112, 2024. doi: 10.1109/ieeedata.2024.3482283.
- C. Juvekar, V. Vaikuntanathan, and A. Chandrakasan. GAZELLE: A low latency framework for secure neural network inference. In *27th USENIX Security Symposium (USENIX Security 18)*, pages 1651–1669, 2018. URL <https://www.usenix.org/conference/usenixsecurity18/presentation/juvekar>.
- M. Kalkatawi, A. Magana-Mora, B. Jankovic, and V. B. Bajic. Deepgsr: an optimized deep-learning structure for the recognition of genomic signals and regions. *Bioinformatics*, 35(7):1125–1132, 2019. doi: 10.1093/bioinformatics/bty752.
- T.-T. Kuo, X. Jiang, H. Tang, X. Wang, T. Bath, D. Bu, L. Wang, A. Harmanci, S. Zhang, D. Zhi, H. J. Sofia, and L. Ohno-Machado. iDASH secure genome analysis competition 2018: blockchain genomic data access logging, homomorphic encryption on GWAS, and dna segment searching. *BMC Medical Genomics*, 13(Suppl 7):98, 2020. doi: 10.1186/s12920-020-0715-0.
- B. Li and D. Micciancio. On the security of homomorphic encryption on approximate numbers. In *Advances in Cryptology – EUROCRYPT 2021*, pages 648–677, 2021. doi: 10.1007/978-3-030-77870-5_23. URL <https://eprint.iacr.org/2020/1533>.
- Z. Li, K. Yang, J. Tan, W.-j. Lu, H. Wu, X. Wang, Y. Yu, D. Zhao, Y. Zheng, M. Guo, and J. Leng. Nimbus: Secure and efficient two-party inference for transformers. In *Advances in Neural Information Processing Systems*, volume 37, pages 21572–21600, 2024. doi: 10.52202/079017-0680.
- W.-j. Lu, Z. Huang, Z. Gu, J. Li, J. Liu, C. Hong, K. Ren, T. Wei, and W. Chen. Bumblebee: Secure two-party inference framework for large transformers. In *Network and Distributed System Security Symposium*, 2025. doi: 10.14722/ndss.2025.230057.
- J. Moon, D. Yoo, X. Jiang, and M. Kim. THOR: Secure transformer inference with homomorphic encryption. In *Proceedings of the 2025 ACM SIGSAC Conference on Computer and Communications Security*, pages 3765–3779, 2025. doi: 10.1145/3719027.3765150. URL <https://eprint.iacr.org/2024/1881>.
- S. Ouaari, J. Kreuer, and N. Pfeifer. How private are DNA embeddings? inverting foundation model representations of genomic sequences, 2026. URL <https://arxiv.org/abs/2603.06950>.

- 802 Q. Pang, J. Zhu, H. Möllering, W. Zheng, and T. Schneider. BOLT: Privacy-preserving, accurate and efficient
803 inference for transformers. In *2024 IEEE Symposium on Security and Privacy (SP)*, pages 4753–4771, 2024. doi:
804 10.1109/SP54263.2024.00130.
- 805 D. Park, E. Lee, and J.-W. Lee. Powerformer: Efficient and high-accuracy privacy-preserving language model
806 with homomorphic encryption. In *Proceedings of the 63rd Annual Meeting of the Association for Computa-
807 tional Linguistics (Volume 1: Long Papers)*, pages 11090–11111, 2025. doi: 10.18653/v1/2025.acl-long.543. URL
808 <https://aclanthology.org/2025.acl-long.543/>.
- 809 L. Yang, S. Wang, L. Yang, K. Wu, and L. Dai. Secure transformer-based neural network inference for protein sequence
810 classification. Cryptology ePrint Archive, Paper 2024/1851, 2024. URL <https://eprint.iacr.org/2024/1851>.
811 First-place solution, iDASH 2024 Track 1.
- 812 Zama. Concrete ML: Privacy-preserving machine learning with fully homomorphic encryption. Software repository,
813 2022a. URL <https://github.com/zama-ai/concrete-ml>.
- 814 Zama. TFHE-rs: A pure rust implementation of the TFHE scheme for boolean and integer arithmetic. Software
815 repository, 2022b. URL <https://github.com/zama-ai/tfhe-rs>.
- 816 D. Zhang, W. Zhang, Y. Zhao, J. Zhang, B. He, C. Qin, and J. Yao. Dnagpt: A generalized pre-trained tool for versatile
817 dna sequence analysis tasks. *arXiv preprint arXiv:2307.05628*, 2023. URL <https://arxiv.org/abs/2307.05628>.
- 818 J. Zhang, X. Yang, L. He, K. Chen, W.-j. Lu, Y. Wang, X. Hou, J. Liu, K. Ren, and X. Yang. Secure transformer
819 inference made non-interactive. In *Network and Distributed System Security Symposium (NDSS)*, 2025. doi: 10.
820 14722/ndss.2025.230868. Also IACR ePrint 2024/136.
- 821 L. Zhang, X. Wang, J. J. Sim, Z. Huang, J. Zhong, H. Wang, P. Duan, and K.-Y. Lam. MOAI: Module-optimizing ar-
822 chitecture for non-interactive secure transformer inference. In *The Fourteenth International Conference on Learning
823 Representations*, 2026. URL <https://openreview.net/forum?id=qJn4HtTzhH>.
- 824 Z. Zhou, Y. Ji, W. Li, P. Dutta, R. V. Davuluri, and H. Liu. Dnabert-2: Efficient foundation model and benchmark
825 for multi-species genomes. In *The Twelfth International Conference on Learning Representations*, 2024. URL
826 <https://openreview.net/forum?id=oMLQB4EZE1>.