

The released model, reproduced locally

Signal recognition and abundance regression use the released fine-tuned heads. The promoter and splice-site rows use a locally fine-tuned head, as none was released.

