


The task that was measured is the longest one

Token counts follow from the task: base pairs divided by six, plus task specials. The twelve-block figures are projections.



Excluded: mRNA abundance regression at 1,755 tokens needs 220 ciphertext groups and roughly 48,000 causal score tiles — a different problem, not a longer one.